

The Strong Lottery Ticket Hypothesis and the Random Subset Sum Problem



Aalto University

Francesco d'Amore

Based on joint work with [A. da Cunha](#) and [E. Natale](#) [NeurIPS 2023]


Aalto Theory Seminar

15 November 2023

Artificial neural networks are large

Usually ranging from **millions** to **hundreds of billions** parameters

- RESNET-50: > 20 millions parameters [He et al. 2015]
- BERT: > 100 millions parameters [Devlin et al. 2018]
- GPT-3: > 100 billions parameters [Brown et al. 2020]




Training is expensive

- Resource intensive
- Good results
- Resulting network still large




Training is expensive

- Resource intensive
- Good results
- Resulting network still large
- Removing edges (**pruning**) works well




Training is expensive

- Resource intensive
- Good results
- Resulting network still large
- Removing edges (**pruning**) works well
- Pruning $\sim 60 - 80\%$ of the edges can lead to better accuracies [Diffenderfer and Kailkhura 2021]




Training is expensive

- Resource intensive
- Good results
- Resulting network still large
- Removing edges ([pruning](#)) works well
- Pruning $\sim 60 - 80\%$ of the edges can lead to better accuracies [\[Diffenderfer and Kailkhura 2021\]](#)
- Pruning $\sim 99\%$ of the edges can perform well [\[Hoefler et al. 2021\]](#)




What if we train the tiny one?

- Maybe, we can avoid the effort of dense training




What if we train the tiny one?

- Maybe, we can avoid the effort of dense training
- Let's test the subnetwork by **retraining** it




What if we train the tiny one?

- Maybe, we can avoid the effort of dense training
- Let's test the subnetwork by **retraining** it
 - Reinitialize




What if we train the tiny one?


- Maybe, we can avoid the effort of dense training
- Let's test the subnetwork by **retraining** it
 - Reinitialize
 - Train



train




prune




train

What if we train the tiny one?


- Maybe, we can avoid the effort of dense training
- Let's test the subnetwork by **retraining** it
 - Reinitialize
 - Train
 - **Bad accuracies**




train



prune




train



Not reinitialization, rewind instead

- Starting from a random point might be too much




Not reinitialization, rewind instead

- Starting from a random point might be too much

[Frankle and Carbin ICLR '19]

- Rewind** instead




Not reinitialization, rewind instead

- Starting from a random point might be too much

[Frankle and Carbin ICLR '19]

- **Rewind** instead
- Training is efficient: 10%-20% of the original size




Not reinitialization, rewind instead

- Starting from a random point might be too much


[Frankle and Carbin ICLR '19]

- **Rewind** instead
- Training is efficient: 10%-20% of the original size
- Similar accuracy




Lottery tickets


- What does it mean?




train



prune




train




Lottery tickets


- What does it mean?
- This is **not** a **good** algorithm




train



prune




train




Lottery tickets


- What does it mean?
- This is **not** a **good algorithm**
- **Existential result**
 - Training is about topology + initialization



???




train



The Lottery Ticket Hypothesis (LTH)


[Frankle and Carbin ICLR '19]: winning lottery tickets always exists



The Lottery Ticket Hypothesis (LTH)

[Frankle and Carbin ICLR '19]: winning lottery tickets always exists

Conjecture: *every randomly-initialized dense network g contains a subnetwork f that matches the test accuracy of g once trained for at most the same number of iterations*




The Lottery Ticket Hypothesis (LTH)

[Frankle and Carbin ICLR '19]: winning lottery tickets always exists

Conjecture: *every randomly-initialized dense network g contains a subnetwork f that matches the test accuracy of g once trained for at most the same number of iterations*

Lot of subsequent work ...



The *Strong* Lottery Ticket Hypothesis (SLTH)

Intuition

- The considered networks are **very large**, and **random**

The *Strong* Lottery Ticket Hypothesis (SLTH)


Intuition

- The considered networks are **very large**, and **random**
- They might already contain **good subnetworks** *from scratch*!

The *Strong* Lottery Ticket Hypothesis (SLTH)

Intuition


- The considered networks are **very large**, and **random**
- They might already contain **good subnetworks** *from scratch*!



The *Strong* Lottery Ticket Hypothesis (SLTH)

Intuition

- The considered networks are **very large**, and **random**
- They might already contain **good subnetworks** *from scratch*!




The *Strong* Lottery Ticket Hypothesis (SLTH)


Intuition

- The considered networks are **very large**, and **random**
- They might already contain **good subnetworks** *from scratch*!

Learn by pruning



prune




The *Strong* Lottery Ticket Hypothesis (SLTH)

Intuition


- The considered networks are **very large**, and **random**
- They might already contain **good subnetworks** *from scratch*!

Learn by pruning



prune

Strong winning lottery ticket



The *Strong* Lottery Ticket Hypothesis (SLTH)

Conjecture: *every randomly-initialized and sufficiently large network g contains a subnetwork f that matches the post-training test accuracy of g even without any training*

The *Strong* Lottery Ticket Hypothesis (SLTH)

Conjecture: *every randomly-initialized and sufficiently large network g contains a subnetwork f that matches the post-training test accuracy of g even without any training*

[Zhou et al. NeurIPS '19] proposes a way to find f : prune weights according to some probability learned through stochastic gradient descent

The *Strong* Lottery Ticket Hypothesis (SLTH)

Conjecture: *every randomly-initialized and sufficiently large network g contains a subnetwork f that matches the post-training test accuracy of g even without any training*

[Zhou et al. NeurIPS '19] proposes a way to find f : prune weights according to some probability learned through stochastic gradient descent

- Decent accuracy

The *Strong* Lottery Ticket Hypothesis (SLTH)

Conjecture: *every randomly-initialized and sufficiently large network g contains a subnetwork f that matches the post-training test accuracy of g even without any training*

[Zhou et al. NeurIPS '19] proposes a way to find f : prune weights according to some probability learned through stochastic gradient descent

- Decent accuracy

[Ramanujan et al. CVPR '20] improves on it: random ResNet-50 pruned to match ResNet-34 on ImageNet

The *Strong* Lottery Ticket Hypothesis (SLTH)

Conjecture: *every randomly-initialized and sufficiently large network g contains a subnetwork f that matches the post-training test accuracy of g even without any training*

[Zhou et al. NeurIPS '19] proposes a way to find f : prune weights according to some probability learned through stochastic gradient descent

- Decent accuracy

[Ramanujan et al. CVPR '20] improves on it: random ResNet-50 pruned to match ResNet-34 on ImageNet

[Diffenderfer and Kailkhura ICLR '21]: works even with binary weights!

Do we have a theorem?

Target result: *Given a network g with random weights, with high probability, it is possible to prune g to approximate any sufficiently smaller network f*

Do we have a theorem?

Target result: *Given a network g with random weights, with high probability, it is possible to prune g to approximate any sufficiently smaller network f*

Target result (equivalent): *Let \mathcal{F} be the class of neural networks with a given size. If a network g with random weights is sufficiently large, then, with high probability, it is possible to prune g to approximate any network in \mathcal{F}*

Do we have a theorem?

Target result: *Given a network g with random weights, with high probability, it is possible to prune g to approximate any sufficiently smaller network f*

Target result (equivalent): *Let \mathcal{F} be the class of neural networks with a given size. If a network g with random weights is sufficiently large, then, with high probability, it is possible to prune g to approximate any network in \mathcal{F}*

- Size: parameter count and depth

Do we have a theorem?

Target result: *Given a network g with random weights, with high probability, it is possible to prune g to approximate any sufficiently smaller network f*

Target result (equivalent): *Let \mathcal{F} be the class of neural networks with a given size. If a network g with random weights is sufficiently large, then, with high probability, it is possible to prune g to approximate any network in \mathcal{F}*

- Size: parameter count and depth
- With high probability: $1 - \delta$ for any given $\delta > 0$

Do we have a theorem?

Target result: *Given a network g with random weights, with high probability, it is possible to prune g to approximate any sufficiently smaller network f*

Target result (equivalent): *Let \mathcal{F} be the class of neural networks with a given size. If a network g with random weights is sufficiently large, then, with high probability, it is possible to prune g to approximate any network in \mathcal{F}*

- Size: parameter count and depth
- With high probability: $1 - \delta$ for any given $\delta > 0$
- Approximation: distance w.r.t. some metric is ε for any given $\varepsilon > 0$

Overview of the (theoretical) results

SLTH holds for:

- [Malach et al. ICML '20]: polynomially overparameterized dense networks with ReLU activation functions

Overview of the (theoretical) results

SLTH holds for:

- [Malach et al. ICML '20]: polynomially overparameterized dense networks with ReLU activation functions
- [Pensia et al. NeurIPS '20]: logarithmically overparameterized dense networks with ReLU activation functions

Overview of the (theoretical) results

SLTH holds for:

- [Malach et al. ICML '20]: polynomially overparameterized dense networks with ReLU activation functions
- [Pensia et al. NeurIPS '20]: logarithmically overparameterized dense networks with ReLU activation functions
- [Diffenderfer and Kailkhura ICLR '21]: polynomially overparameterized binary dense networks

Overview of the (theoretical) results

SLTH holds for:

- [Malach et al. ICML '20]: polynomially overparameterized dense networks with ReLU activation functions
- [Pensia et al. NeurIPS '20]: logarithmically overparameterized dense networks with ReLU activation functions
- [Diffenderfer and Kailkhura ICLR '21]: polynomially overparameterized binary dense networks
- [Sreenivasan et al. AISTAT '22]: polylogarithmically overparameterized binary dense networks

Overview of the (theoretical) results

SLTH holds for:

- [Malach et al. ICML '20]: **polynomially overparameterized** dense networks with ReLU activation functions
- [Pensia et al. NeurIPS '20]: **logarithmically overparameterized** dense networks with ReLU activation functions
- [Diffenderfer and Kailkhura ICLR '21]: **polynomially overparameterized** binary dense networks
- [Sreenivasan et al. AISTAT '22]: **polylogarithmically overparameterized** binary dense networks
- [da Cunha et al. ICLR '22]: **logarithmically overparameterized** convolutional neural networks (CNNs) with ReLU activation functions and non-negative inputs

Overview of the (theoretical) results

SLTH holds for:

- [Malach et al. ICML '20]: **polynomially overparameterized** dense networks with ReLU activation functions
- [Pensia et al. NeurIPS '20]: **logarithmically overparameterized** dense networks with ReLU activation functions
- [Diffenderfer and Kailkhura ICLR '21]: **polynomially overparameterized** binary dense networks
- [Sreenivasan et al. AISTAT '22]: **polylogarithmically overparameterized** binary dense networks
- [da Cunha et al. ICLR '22]: **logarithmically overparameterized** convolutional neural networks (CNNs) with ReLU activation functions and non-negative inputs
- [Burkholz NeurIPS, ICML '22]: **logarithmically overparameterized** dense networks, CNNs, and residual architectures with a wider class of activation functions and less depth overhead

Overview of the (theoretical) results

SLTH holds for:

- [Malach et al. ICML '20]: **polynomially overparameterized** dense networks with ReLU activation functions
- [Pensia et al. NeurIPS '20]: **logarithmically overparameterized** dense networks with ReLU activation functions
- [Diffenderfer and Kailkhura ICLR '21]: **polynomially overparameterized** binary dense networks
- [Sreenivasan et al. AISTAT '22]: **polylogarithmically overparameterized** binary dense networks
- [da Cunha et al. ICLR '22]: **logarithmically overparameterized** convolutional neural networks (CNNs) with ReLU activation functions and non-negative inputs
- [Burkholz NeurIPS, ICML '22]: **logarithmically overparameterized** dense networks, CNNs, and residual architectures with a wider class of activation functions and less depth overhead
- [Ferbach et al. ICLR '22]: **logarithmically overparameterized** equivariant networks with ReLU activation functions

Review: SLTH in dense networks


Dense network: $f(\mathbf{x}) = \mathbf{W}_\ell \sigma(\mathbf{W}_{\ell-1} \dots \sigma(\mathbf{W}_1 \mathbf{x}))$

- $\mathbf{x} \in \mathbb{R}^{d_0}$, $\mathbf{W}_i \in \mathbb{R}^{d_{i-1} \times d_i}$
- $\sigma(x) = \max(0, x)$ (ReLU)

Review: SLTH in dense networks

Dense network: $f(\mathbf{x}) = \mathbf{W}_\ell \sigma(\mathbf{W}_{\ell-1} \dots \sigma(\mathbf{W}_1 \mathbf{x}))$


- $\mathbf{x} \in \mathbb{R}^{d_0}$, $\mathbf{W}_i \in \mathbb{R}^{d_{i-1} \times d_i}$
- $\sigma(x) = \max(0, x)$ (ReLU)



Review: SLTH in dense networks

Dense network: $f(\mathbf{x}) = \mathbf{W}_\ell \sigma(\mathbf{W}_{\ell-1} \dots \sigma(\mathbf{W}_1 \mathbf{x}))$

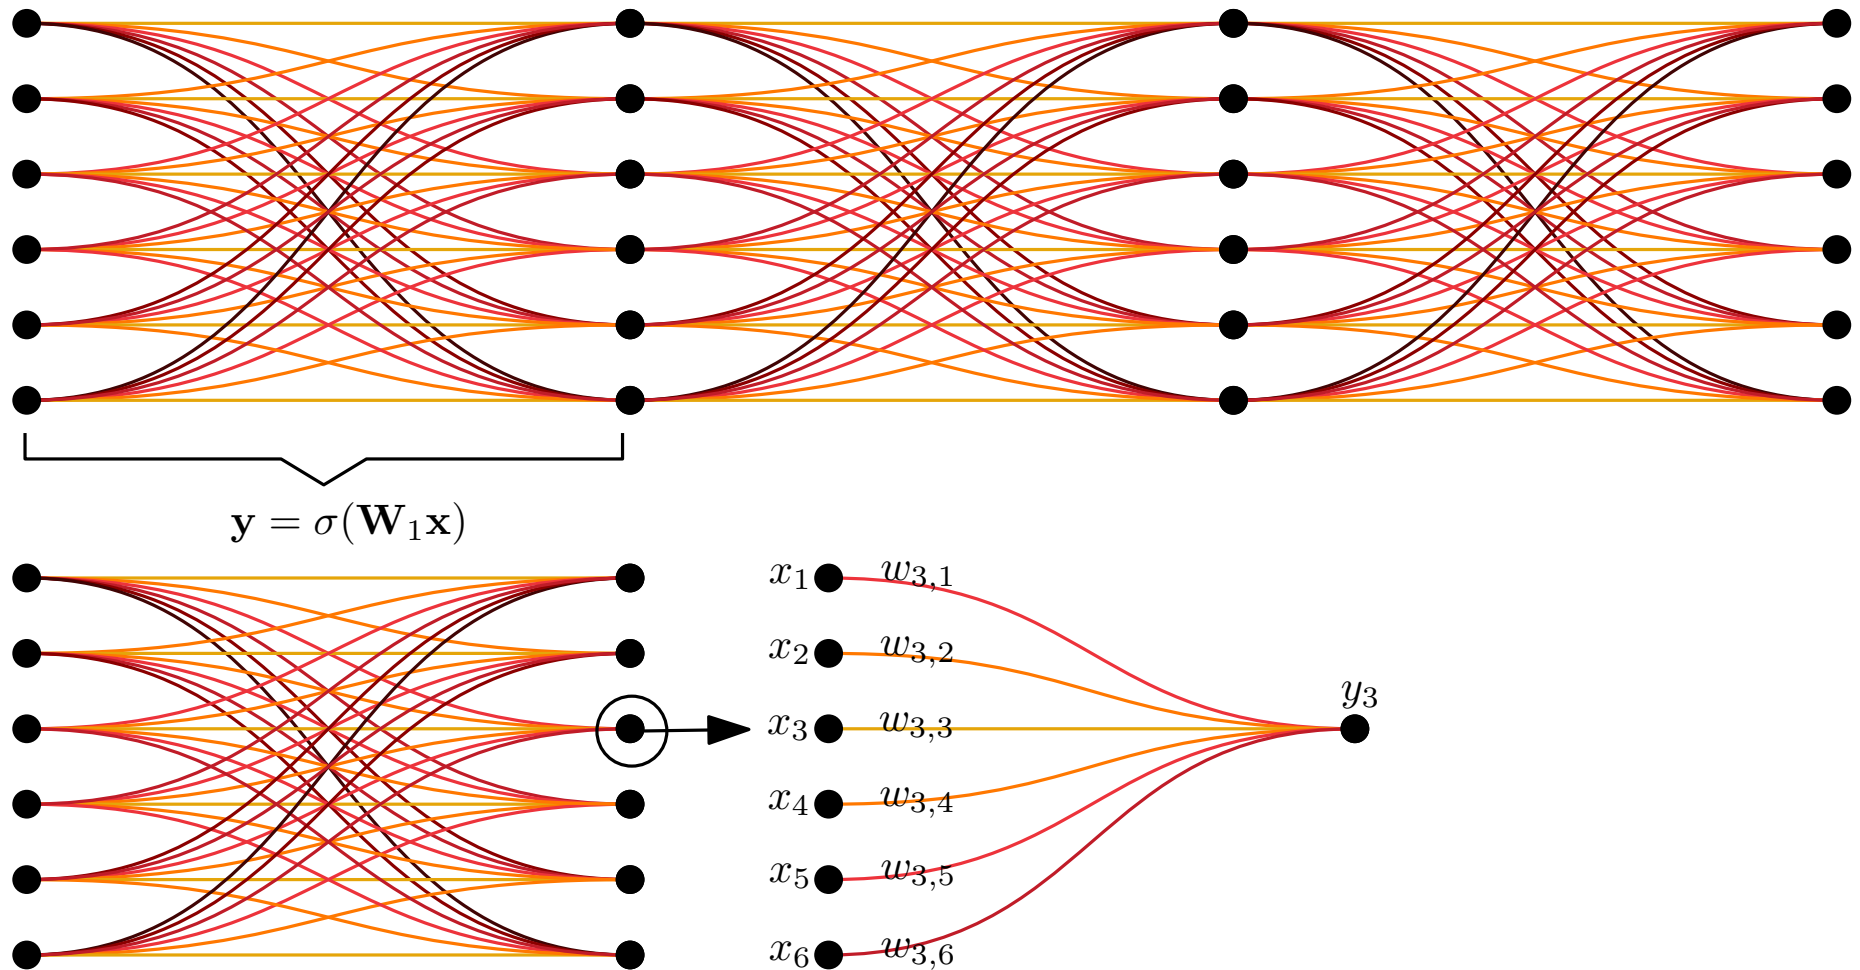
- $\mathbf{x} \in \mathbb{R}^{d_0}$, $\mathbf{W}_i \in \mathbb{R}^{d_{i-1} \times d_i}$
- $\sigma(x) = \max(0, x)$ (ReLU)



Review: SLTH in dense networks

Dense network: $f(\mathbf{x}) = \mathbf{W}_\ell \sigma(\mathbf{W}_{\ell-1} \dots \sigma(\mathbf{W}_1 \mathbf{x}))$

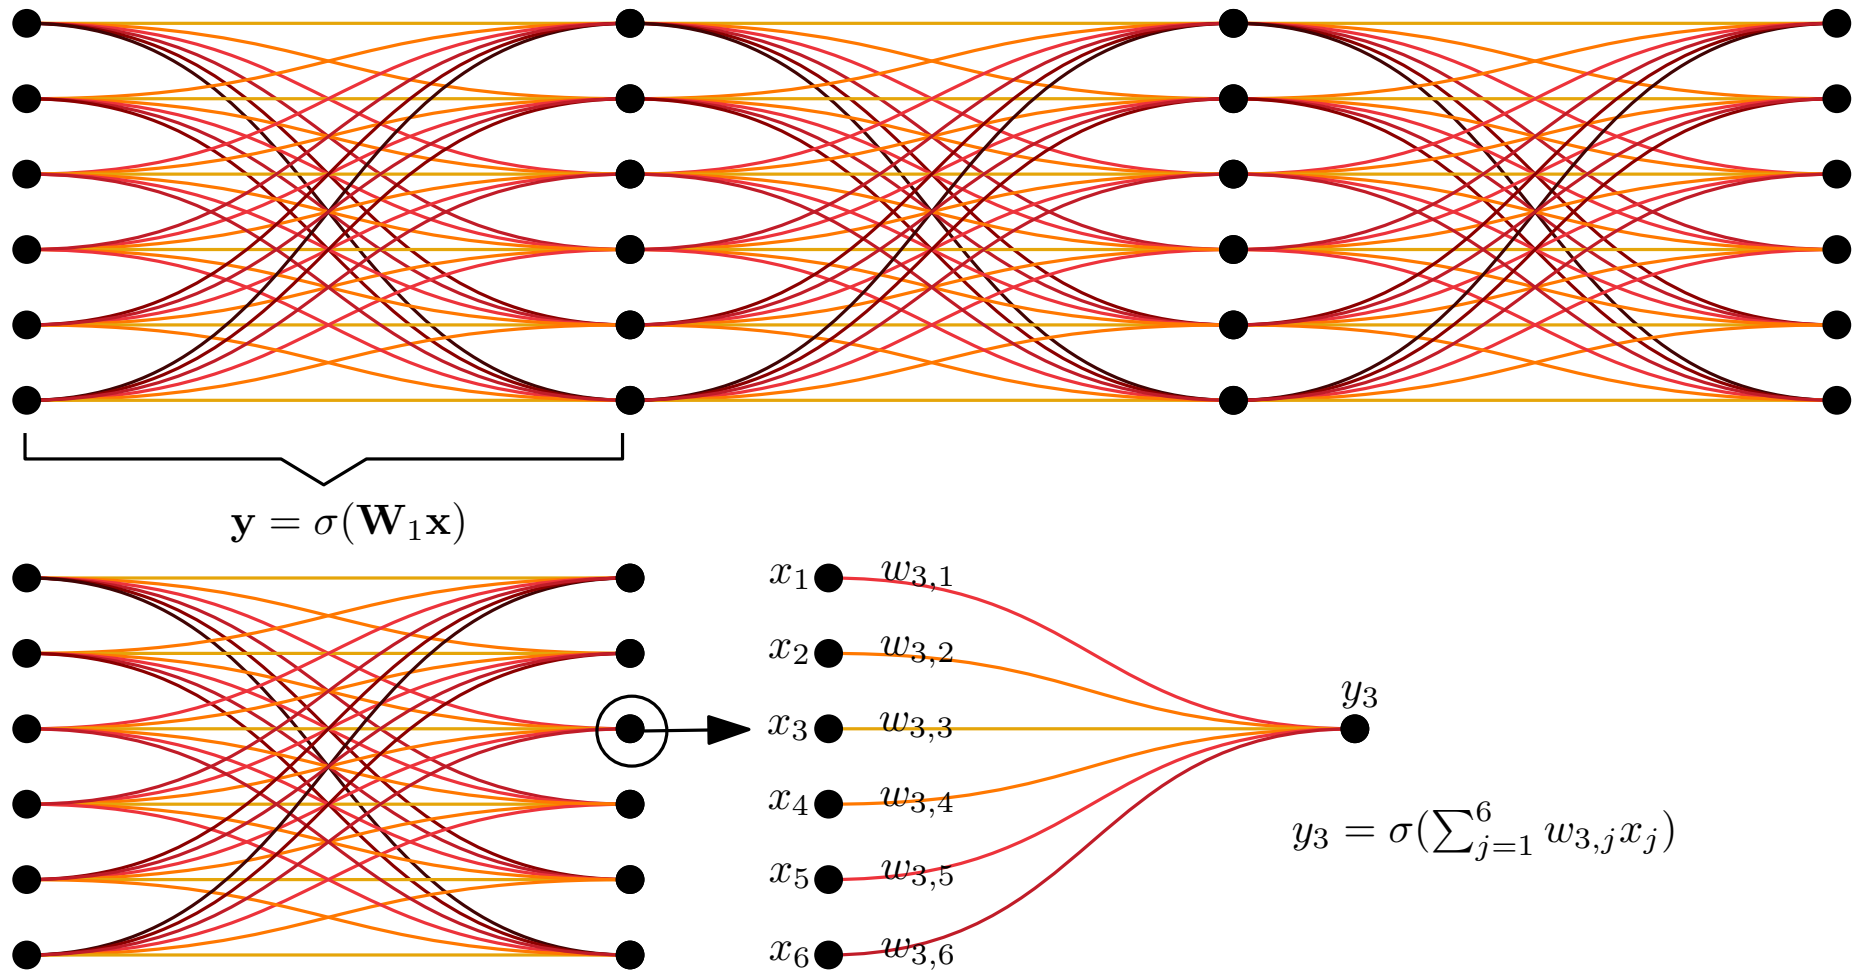
- $\mathbf{x} \in \mathbb{R}^{d_0}$, $\mathbf{W}_i \in \mathbb{R}^{d_{i-1} \times d_i}$
- $\sigma(x) = \max(0, x)$ (ReLU)



Review: SLTH in dense networks


Dense network: $f(\mathbf{x}) = \mathbf{W}_\ell \sigma(\mathbf{W}_{\ell-1} \dots \sigma(\mathbf{W}_1 \mathbf{x}))$

- $\mathbf{x} \in \mathbb{R}^{d_0}$, $\mathbf{W}_i \in \mathbb{R}^{d_{i-1} \times d_i}$
- $\sigma(x) = \max(0, x)$ (ReLU)




Review: SLTH in dense networks

- First target: approx $y = wx$ within error ε (no ReLU, one edge only)



Review: SLTH in dense networks


- First target: approx $y = wx$ within error ε (no ReLU, one edge only)



- Original approach


add **intermediate layer**, sample

$w_i \sim \text{Unif}[-1, 1]$ until getting $w \pm \varepsilon$



Review: SLTH in dense networks


- First target: approx $y = wx$ within error ε (no ReLU, one edge only)



- Original approach

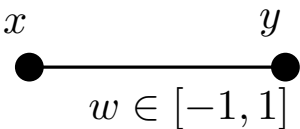
add **intermediate layer**, sample


$w_i \sim \text{Unif}[-1, 1]$ until getting $w \pm \varepsilon$




roughly $1/\varepsilon$ samples

Review: SLTH in dense networks

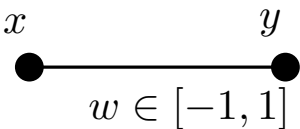
- First target: approx $y = wx$ within error ε (no ReLU, one edge only) 
 $w \in [-1, 1]$
- Original approach
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ until getting $w \pm \varepsilon$
- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**




roughly $1/\varepsilon$ samples




Review: SLTH in dense networks

- First target: approx $y = wx$ within error ε (no ReLU, one edge only) 
 $w \in [-1, 1]$
- Original approach
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ until getting $w \pm \varepsilon$
- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**

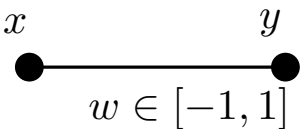



roughly $1/\varepsilon$ samples




How many?

Review: SLTH in dense networks

- First target: approx $y = wx$ within error ε (no ReLU, one edge only) 
- Original approach
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ until getting $w \pm \varepsilon$
- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**



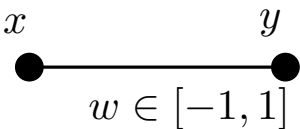
roughly $1/\varepsilon$ samples




How many?


Theorem [Lueker 1998; da Cunha et al. ESA '23]: Let $x_1, \dots, x_n \in [-1, 1]$ be i.i.d. uniform random variables. Given any error parameter $\varepsilon > 0$, there exists a constant $C > 0$ such that if $n \geq C \log 1/\varepsilon$ then, with probability $1 - \exp[-(n - C \log 1/\varepsilon)^2 / 4n]$, for each $z \in [-1, 1]$ there exists a subset $S \subseteq [n]$ such that $|z - \sum_{i \in S} x_i| < 2\varepsilon$

Review: SLTH in dense networks

- First target: approx $y = wx$ within error ε (no ReLU, one edge only) 
- Original approach
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ until getting $w \pm \varepsilon$
- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**



roughly $1/\varepsilon$ samples




How many?

Theorem [Lueker 1998; da Cunha et al. ESA '23]: Let $x_1, \dots, x_n \in [-1, 1]$ be i.i.d. uniform random variables. Given any error parameter $\varepsilon > 0$, there exists a constant $C > 0$ such that if $n \geq C \log 1/\varepsilon$ then, with probability $1 - \exp[-(n - C \log 1/\varepsilon)^2 / 4n]$, for each $z \in [-1, 1]$ there exists a subset $S \subseteq [n]$ such that $|z - \sum_{i \in S} x_i| < 2\varepsilon$

works for all densities $h(x) = pf(x) + (1 - p)g(x)$, where f is “uniform”


RSS approach

- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**



RSS approach


- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**



$$n \geq C \log 1/\varepsilon \implies \exists S \subseteq [n] : \left| w - \sum_{i \in S} w_i \right| < 2\varepsilon$$

RSS approach

- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**




$$n \geq C \log 1/\varepsilon \implies \exists S \subseteq [n] : \left| w - \sum_{i \in S} w_i \right| < 2\varepsilon$$

$$\implies \left| wx - \sum_{i \in S} w_i x \right| \leq |x| \left| w - \sum_{i \in S} w_i \right| < 2\varepsilon |x|$$

RSS approach

- Random subset sum (RSS) approach:


add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**



$$n \geq C \log 1/\varepsilon \implies \exists S \subseteq [n] : |w - \sum_{i \in S} w_i| < 2\varepsilon$$

$$\implies |wx - \sum_{i \in S} w_i x| \leq |x| |w - \sum_{i \in S} w_i| < 2\varepsilon |x|$$


- Completely random initialization + ReLU (**non-linearity**):



ReLU:
 $\sigma(x) = \max(0, x)$

RSS approach


- Random subset sum (RSS) approach:
add **intermediate layer**, sample $w_i \sim \text{Unif}[-1, 1]$ and find a **good subset**



$$n \geq C \log 1/\varepsilon \implies \exists S \subseteq [n] : |w - \sum_{i \in S} w_i| < 2\varepsilon$$

$$\implies |wx - \sum_{i \in S} w_i x| \leq |x| |w - \sum_{i \in S} w_i| < 2\varepsilon |x|$$

- Completely random initialization + ReLU (**non-linearity**):
how to deal with non-linearity?



ReLU:
 $\sigma(x) = \max(0, x)$


Exploiting properties of the ReLU

- Completely random initialization + ReLU (**non-linearity**)

Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$




Exploiting properties of the ReLU

- Completely random initialization + ReLU (non-linearity)

Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:
 $\sigma(x) = \max(0, x)$




Exploiting properties of the ReLU

- Completely random initialization + ReLU (non-linearity)

Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$




Exploiting properties of the ReLU

- Completely random initialization + ReLU (**non-linearity**)

Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$




Exploiting properties of the ReLU

- Completely random initialization + ReLU (**non-linearity**)


Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$



- How? Wlog, assume $w \geq 0$




Exploiting properties of the ReLU

- Completely random initialization + ReLU (**non-linearity**)


Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$



- How? Wlog, assume $w \geq 0$




Exploiting properties of the ReLU

- Completely random initialization + ReLU (**non-linearity**)


Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$



- How? Wlog, assume $w \geq 0$




Exploiting properties of the ReLU

- Completely random initialization + ReLU (**non-linearity**)


Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$



- How? Wlog, assume $w \geq 0$



$$\left| \sigma(w x) - \sum_{i \in [n]} b_i \sigma(a_i x) \right| = \left| \sigma(w x) - \sum_{i \in [n]} b_i a_i^+ \sigma(x) \right|$$

if $x \leq 0$, easy


Exploiting properties of the ReLU

- Completely random initialization + ReLU (**non-linearity**)


Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:

$$\sigma(x) = \max(0, x)$$



- How? Wlog, assume $w \geq 0$




Exploiting properties of the ReLU


- Completely random initialization + ReLU (**non-linearity**)

Property of ReLU: $w x = \sigma(w x) - \sigma(-w x)$

ReLU:
 $\sigma(x) = \max(0, x)$



- How? Wlog, assume $w \geq 0$




if $x \leq 0$, easy


if $x > 0$, then RSS holds!

$n \geq C \log 1/\varepsilon$

Putting everything together




Putting everything together




prune only the right layer: **reuse** the **left layer**


Putting everything together




prune only the right layer: **reuse** the **left** layer




Putting everything together




prune only the right layer: **reuse** the **left** layer




Putting everything together




prune only the right layer: **reuse** the **left** layer



More layers together




More layers together




$$n \geq C \log \ell d^2 / \varepsilon$$

$\ell = \#$ original layers

$$\Rightarrow \|y - z\| \leq 2\varepsilon$$




More layers together



$$n \geq C \log \ell d^2 / \varepsilon$$


$\ell = \#$ original layers

$$\implies \|\mathbf{y} - \mathbf{z}\| \leq 2\varepsilon$$




Unstructured pruning

- Removed edges can be everywhere



Unstructured pruning


- Removed edges can be everywhere



- No structure usually **implies slower processes**

Unstructured pruning


- Removed edges can be everywhere



- No structure usually **implies slower processes**
 - difficulty encoding unstructured sparsity

Unstructured pruning


- Removed edges can be everywhere



- No structure usually **implies slower processes**
 - difficulty encoding unstructured sparsity
 - accessing data is more time consuming than processing

Unstructured pruning


- Removed edges can be everywhere



- No structure usually **implies slower processes**
 - difficulty encoding unstructured sparsity
 - accessing data is more time consuming than processing
 - the processor register allows parallel operations for blocks of memory


Unstructured pruning

- Removed edges can be everywhere




- No structure usually **implies slower processes**
 - difficulty encoding unstructured sparsity
 - accessing data is more time consuming than processing
 - the processor register allows parallel operations for blocks of memory
- [Malach et al. ICML '20]: pruning neurons alone requires **exponential** overparameterization


Structured pruning




Structured pruning



Structured pruning




Structured pruning




- Removing entire neurons from the middle layer!

Structured pruning



- Removing entire neurons from the middle layer!


Structured pruning



- Removing entire neurons from the middle layer!
- removes columns!

$$\begin{bmatrix} 0 & v_{1,2} & 0 & \dots & 0 & v_{i,1} & 0 & \dots \\ 0 & v_{2,2} & 0 & \dots & 0 & v_{i,2} & 0 & \dots \\ 0 & v_{3,2} & 0 & \dots & 0 & v_{i,2} & 0 & \dots \end{bmatrix} \cdot \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_{3n} \end{bmatrix}$$


Structured pruning



- Removing entire neurons from the middle layer!
- removes columns!
- The one-dimensional RSS result does not work
- leads to **exponential** bounds

$$\begin{bmatrix} 0 & v_{1,2} & 0 & \dots & 0 & v_{i,1} & 0 & \dots \\ 0 & v_{2,2} & 0 & \dots & 0 & v_{i,2} & 0 & \dots \\ 0 & v_{3,2} & 0 & \dots & 0 & v_{i,2} & 0 & \dots \end{bmatrix} \cdot \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_{3n} \end{bmatrix}$$

Structured pruning




- Removing entire neurons from the middle layer!
- removes columns!
- The one-dimensional RSS result does not work
- leads to **exponential** bounds
- A **multidimensional** RSS result is required

$$\begin{bmatrix} 0 & v_{1,2} & 0 & \dots & 0 & v_{i,1} & 0 & \dots \\ 0 & v_{2,2} & 0 & \dots & 0 & v_{i,2} & 0 & \dots \\ 0 & v_{3,2} & 0 & \dots & 0 & v_{i,2} & 0 & \dots \end{bmatrix} \cdot \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_{3n} \end{bmatrix}$$

The multidimensional RSS problem

- Natural generalization




The multidimensional RSS problem

- Natural generalization

Input:

- Sequence of n i.i.d. random vectors X_1, \dots, X_n




The multidimensional RSS problem

- Natural generalization

Input:

- **Sequence** of n i.i.d. **random vectors** X_1, \dots, X_n
- **Target** vector $\mathbf{z} \in [-1, +1]^d$




The multidimensional RSS problem

- Natural generalization

Input:

- **Sequence** of n i.i.d. **random vectors** X_1, \dots, X_n
- **Target** vector $\mathbf{z} \in [-1, +1]^d$
- **Error** parameter $\varepsilon > 0$




The multidimensional RSS problem

- Natural generalization

Input:

- **Sequence** of n i.i.d. **random vectors** X_1, \dots, X_n
- **Target** vector $\mathbf{z} \in [-1, +1]^d$
- **Error** parameter $\varepsilon > 0$




Question:

- Estimate n such that, with high probability, a **subset** $S \subseteq [n]$ exists with
$$\|\mathbf{z} - \sum_{i \in S} X_i\|_{\infty} \leq 2\varepsilon$$


MRSS in expectation

- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$




MRSS in expectation

- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$




MRSS in expectation

- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$
- 2^n possible subsets



MRSS in expectation


- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$
- 2^n possible subsets
- Target $\mathbf{z} \in [-1, 1]^d$



MRSS in expectation


- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$
- 2^n possible subsets
- Target $\mathbf{z} \in [-1, 1]^d$

Upper bound



MRSS in expectation

- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$
- 2^n possible subsets
- Target $\mathbf{z} \in [-1, 1]^d$




Upper bound

- If subset size $k = \frac{n}{2}$, possible subsets: $\binom{n}{n/2} \geq 2^{n/2}$

MRSS in expectation

- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$
- 2^n possible subsets
- Target $\mathbf{z} \in [-1, 1]^d$




Upper bound

- If subset size $k = \frac{n}{2}$, possible subsets: $\binom{n}{n/2} \geq 2^{n/2}$
- Each subset $S \subseteq [n]$, $|S| = \frac{n}{2}$, gives a Gaussian $Y_S \sim \mathcal{N}(\mathbf{0}, \frac{n}{2}I_d)$

MRSS in expectation

- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$
- 2^n possible subsets
- Target $\mathbf{z} \in [-1, 1]^d$




Upper bound

- If subset size $k = \frac{n}{2}$, possible subsets: $\binom{n}{n/2} \geq 2^{n/2}$
- Each subset $S \subseteq [n]$, $|S| = \frac{n}{2}$, gives a Gaussian $Y_S \sim \mathcal{N}(\mathbf{0}, \frac{n}{2}I_d)$
- Probability roughly $(\varepsilon/\sqrt{n/2})^d$ to hit any ε -cube

MRSS in expectation

- Number of ε -cubes: $1/\varepsilon^d = 2^{d \log 1/\varepsilon}$
- **Sequence** of n i.i.d. **random vectors**
 $X_1, \dots, X_n \sim \mathcal{N}(\mathbf{0}, I_d)$
- 2^n possible subsets
- Target $\mathbf{z} \in [-1, 1]^d$



Upper bound

- If subset size $k = \frac{n}{2}$, possible subsets: $\binom{n}{n/2} \geq 2^{n/2}$
- Each subset $S \subseteq [n]$, $|S| = \frac{n}{2}$, gives a Gaussian $Y_S \sim \mathcal{N}(\mathbf{0}, \frac{n}{2}I_d)$
- Probability roughly $(\varepsilon/\sqrt{n/2})^d$ to hit any ε -cube

$$\begin{aligned} \mathbb{E}[\# \text{ subsets approximating any cube}] &\geq 2^{n/2} \cdot \left(\frac{\varepsilon}{\sqrt{n/2}} \right)^d \\ &= 2^{n/2 - d \log 1/\varepsilon - d/2 \log n/2} = 2^{O(n)} \text{ if } n \geq Cd \log 1/\varepsilon \end{aligned}$$

MRSS in expectation

Lower bound

- If subset size k , possible subsets: $\binom{n}{k} \leq (en/k)^k$

MRSS in expectation

Lower bound

- If subset size k , possible subsets: $\binom{n}{k} \leq (en/k)^k$
- Each subset $S \subseteq [n]$, $|S| = k$, gives a Gaussian $Y_S \sim \mathcal{N}(\mathbf{0}, kI_d)$

MRSS in expectation

Lower bound

- If subset size k , possible subsets: $\binom{n}{k} \leq (en/k)^k$
- Each subset $S \subseteq [n]$, $|S| = k$, gives a Gaussian $Y_S \sim \mathcal{N}(\mathbf{0}, kI_d)$
- Probability roughly $(\varepsilon/\sqrt{k})^d$ to hit any ε -cube

MRSS in expectation

Lower bound

- If subset size k , possible subsets: $\binom{n}{k} \leq (en/k)^k$
- Each subset $S \subseteq [n]$, $|S| = k$, gives a Gaussian $Y_S \sim \mathcal{N}(\mathbf{0}, kI_d)$
- Probability roughly $(\varepsilon/\sqrt{k})^d$ to hit any ε -cube

$$\begin{aligned}\mathbb{E} [\# \text{ subsets approximating any cube}] &\leq \sum_{k=1}^n (en/k)^k \cdot \left(\frac{\varepsilon}{\sqrt{k}}\right)^d \\ &= \sum_{k=1}^n 2^{k \log(en/k) - d \log 1/\varepsilon - d/2 \log k} \leq n \cdot 2^{n/2 \log(2e) - d \log 1/\varepsilon - d/2 \log n/2}\end{aligned}$$

MRSS in expectation

Lower bound

- If subset size k , possible subsets: $\binom{n}{k} \leq (en/k)^k$
- Each subset $S \subseteq [n]$, $|S| = k$, gives a Gaussian $Y_S \sim \mathcal{N}(\mathbf{0}, kI_d)$
- Probability roughly $(\varepsilon/\sqrt{k})^d$ to hit any ε -cube

$$\begin{aligned} \mathbb{E} [\# \text{ subsets approximating any cube}] &\leq \sum_{k=1}^n (en/k)^k \cdot \left(\frac{\varepsilon}{\sqrt{k}} \right)^d \\ &= \sum_{k=1}^n 2^{k \log(en/k) - d \log 1/\varepsilon - d/2 \log k} \leq n \cdot 2^{n/2 \log(2e) - d \log 1/\varepsilon - d/2 \log n/2} \\ &< 1 \text{ if } n \leq cd \log 1/\varepsilon \text{ for } c \text{ small enough} \end{aligned}$$

MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the 2nd moment method to derive bounds

MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the 2nd moment method to derive bounds
 - for $S \subseteq [n]$, $Y_S = 1$ if $\sum_{i \in S} X_i$ approximates target z and 0 otherwise

MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the 2nd moment method to derive bounds
 - for $S \subseteq [n]$, $Y_S = 1$ if $\sum_{i \in S} X_i$ approximates target z and 0 otherwise
 - $Z_n = \sum_{S \subseteq [n]} Y_S$ number of subsets approximating target z

MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the 2nd moment method to derive bounds
 - for $S \subseteq [n]$, $Y_S = 1$ if $\sum_{i \in S} X_i$ approximates target z and 0 otherwise
 - $Z_n = \sum_{S \subseteq [n]} Y_S$ number of subsets approximating target z
 - $\mathbb{P}[Z_n \geq 1] \geq (\mathbb{E}[Z_n])^2 / \mathbb{E}[Z_n^2]$

MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the 2nd moment method to derive bounds
 - for $S \subseteq [n]$, $Y_S = 1$ if $\sum_{i \in S} X_i$ approximates target z and 0 otherwise
 - $Z_n = \sum_{S \subseteq [n]} Y_S$ number of subsets approximating target z
 - $\mathbb{P}[Z_n \geq 1] \geq (\mathbb{E}[Z_n])^2 / \mathbb{E}[Z_n^2]$
- **Challenge:** dealing with dependencies to estimate $\mathbb{E}[Z_n^2]$

MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the **2nd moment method** to derive bounds
 - for $S \subseteq [n]$, $Y_S = 1$ if $\sum_{i \in S} X_i$ approximates target z and 0 otherwise
 - $Z_n = \sum_{S \subseteq [n]} Y_S$ **number of subsets** approximating target z
 - $\mathbb{P}[Z_n \geq 1] \geq (\mathbb{E}[Z_n])^2 / \mathbb{E}[Z_n^2]$
- **Challenge:** dealing with **dependencies** to estimate $\mathbb{E}[Z_n^2]$
 - choose only subsets of size αn so that the “average intersection” concentrates around $\alpha^2 n$


MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the **2nd moment method** to derive bounds
 - for $S \subseteq [n]$, $Y_S = 1$ if $\sum_{i \in S} X_i$ approximates target \mathbf{z} and 0 otherwise
 - $Z_n = \sum_{S \subseteq [n]} Y_S$ **number of subsets** approximating target \mathbf{z}
 - $\mathbb{P}[Z_n \geq 1] \geq (\mathbb{E}[Z_n])^2 / \mathbb{E}[Z_n^2]$
- **Challenge:** dealing with **dependencies** to estimate $\mathbb{E}[Z_n^2]$
 - choose only subsets of size αn so that the “average intersection” concentrates around $\alpha^2 n$
- **Result:** $n \geq \text{poly}(d) \log(d/\varepsilon)$ ($\alpha = 1/\sqrt{d}$)


MRSS: current results

- [Borst et al. 2022; Becchetti et al. 2022] use the **2nd moment method** to derive bounds
 - for $S \subseteq [n]$, $Y_S = 1$ if $\sum_{i \in S} X_i$ approximates target \mathbf{z} and 0 otherwise
 - $Z_n = \sum_{S \subseteq [n]} Y_S$ **number of subsets** approximating target \mathbf{z}
 - $\mathbb{P}[Z_n \geq 1] \geq (\mathbb{E}[Z_n])^2 / \mathbb{E}[Z_n^2]$
- **Challenge:** dealing with **dependencies** to estimate $\mathbb{E}[Z_n^2]$
 - choose only subsets of size αn so that the “average intersection” concentrates around $\alpha^2 n$
- **Result:** $n \geq \text{poly}(d) \log(d/\varepsilon)$ ($\alpha = 1/\sqrt{d}$)
- What about approximating all the hypercube $[-1, 1]^d$? The **union bound** is highly **non-optimal**


Apply MRSS for structured pruning




Apply MRSS for structured pruning




Apply MRSS for structured pruning




Apply MRSS for structured pruning



Apply MRSS for structured pruning




Apply MRSS for structured pruning




• $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)

• For simplicity: **no ReLU**




Apply MRSS for structured pruning




• $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)

• For simplicity: **no ReLU**

$$\|x_1 \mathbf{w}_1 - \sum_{i=1}^n x_1 a_{1,i} \mathbf{b}_i\|_{\infty} \leq |x_1| \|\mathbf{w}_1 - \sum_{i=1}^n a_{1,i} \mathbf{b}_i\|_{\infty}$$



Apply MRSS for structured pruning




- $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)


- For simplicity: **no ReLU**

$$\|x_1 \mathbf{w}_1 - \sum_{i=1}^n x_1 a_{1,i} \mathbf{b}_i\|_{\infty} \leq |x_1| \|\mathbf{w}_1 - \sum_{i=1}^n a_{1,i} \mathbf{b}_i\|_{\infty}$$

- **Issue:** dependencies among entries of $a_{1,i} \mathbf{b}_i$!



Apply MRSS for structured pruning



- $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)


- For simplicity: **no ReLU**

$$\|x_1 \mathbf{w}_1 - \sum_{i=1}^n x_1 a_{1,i} \mathbf{b}_i\|_{\infty} \leq |x_1| \|\mathbf{w}_1 - \sum_{i=1}^n a_{1,i} \mathbf{b}_i\|_{\infty}$$


- **Issue:** dependencies among entries of $a_{1,i} \mathbf{b}_i$!

- **Solution:**

$$\text{- for } S \subseteq [n], X_S = \sum_{i \in S} a_{1,i} \mathbf{b}_i$$



Apply MRSS for structured pruning



- $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)

- For simplicity: **no ReLU**


$$\|x_1 \mathbf{w}_1 - \sum_{i=1}^n x_1 a_{1,i} \mathbf{b}_i\|_{\infty} \leq |x_1| \|\mathbf{w}_1 - \sum_{i=1}^n a_{1,i} \mathbf{b}_i\|_{\infty}$$

- **Issue:** dependencies among entries of $a_{1,i} \mathbf{b}_i$!


- **Solution:**

- for $S \subseteq [n]$, $X_S = \sum_{i \in S} a_{1,i} \mathbf{b}_i$

- conditional on $a_{1,i}$ for each $i \in S$, X_S is distributed as $\mathcal{N}(\mathbf{0}, \sum_{i \in S} a_{1,i}^2 \cdot I_d)$



Apply MRSS for structured pruning



- $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)

- For simplicity: **no ReLU**

$$\|x_1 \mathbf{w}_1 - \sum_{i=1}^n x_1 a_{1,i} \mathbf{b}_i\|_{\infty} \leq |x_1| \|\mathbf{w}_1 - \sum_{i=1}^n a_{1,i} \mathbf{b}_i\|_{\infty}$$


- **Issue:** dependencies among entries of $a_{1,i} \mathbf{b}_i$!

- **Solution:**


- for $S \subseteq [n]$, $X_S = \sum_{i \in S} a_{1,i} \mathbf{b}_i$

- conditional on $a_{1,i}$ for each $i \in S$, X_S is distributed as $\mathcal{N}(\mathbf{0}, \sum_{i \in S} a_{1,i}^2 \cdot I_d)$

- $\sum_{i \in S} a_{1,i}^2$ is a Chi-squared distribution: **concentration inequalities!**



Apply MRSS for structured pruning



- $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)

- For simplicity: **no ReLU**

$$\|x_1 \mathbf{w}_1 - \sum_{i=1}^n x_1 a_{1,i} \mathbf{b}_i\|_{\infty} \leq |x_1| \|\mathbf{w}_1 - \sum_{i=1}^n a_{1,i} \mathbf{b}_i\|_{\infty}$$

- **Issue:** dependencies among entries of $a_{1,i} \mathbf{b}_i$!


- **Solution:**

- for $S \subseteq [n]$, $X_S = \sum_{i \in S} a_{1,i} \mathbf{b}_i$


- conditional on $a_{1,i}$ for each $i \in S$, X_S is distributed as $\mathcal{N}(\mathbf{0}, \sum_{i \in S} a_{1,i}^2 \cdot I_d)$

- $\sum_{i \in S} a_{1,i}^2$ is a Chi-squared distribution: **concentration inequalities!**

- *things do not change too much*



Apply MRSS for structured pruning



- $\mathbf{a}_i \sim \mathcal{N}(\mathbf{0}, I_n)$, $\mathbf{b}_i \sim \mathcal{N}(\mathbf{0}, I_d)$ (here, $d = 3$)

- For simplicity: **no ReLU**

$$\|x_1 \mathbf{w}_1 - \sum_{i=1}^n x_1 a_{1,i} \mathbf{b}_i\|_{\infty} \leq |x_1| \|\mathbf{w}_1 - \sum_{i=1}^n a_{1,i} \mathbf{b}_i\|_{\infty}$$

- **Issue:** dependencies among entries of $a_{1,i} \mathbf{b}_i$!

- **Solution:**


- for $S \subseteq [n]$, $X_S = \sum_{i \in S} a_{1,i} \mathbf{b}_i$

- conditional on $a_{1,i}$ for each $i \in S$, X_S is distributed as $\mathcal{N}(\mathbf{0}, \sum_{i \in S} a_{1,i}^2 \cdot I_d)$

- $\sum_{i \in S} a_{1,i}^2$ is a Chi-squared distribution: **concentration inequalities!**

- *things do not change too much*

Result: $n \geq \text{poly}(d) \cdot \text{polylog}(d\ell/\varepsilon)$



Convolutional neural networks (CNNs)

- **Generality**: we actually prove the result in CNNs

Convolutional neural networks (CNNs)

- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:**

$$\begin{array}{|c|c|c|} \hline 1 & 3 & -1 \\ \hline 0 & -2 & 0 \\ \hline 4 & -1 & -1 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline 1 & 3 & 2 & 1 \\ \hline 1 & 0 & 0 & -2 \\ \hline 4 & -2 & 3 & 0 \\ \hline 0 & 1 & 2 & -5 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 19 & -3 \\ \hline 2 & 3 \\ \hline \end{array}$$

Convolutional neural networks (CNNs)

- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:**

$$\begin{array}{|c|c|c|} \hline 1 & 3 & -1 \\ \hline 0 & -2 & 0 \\ \hline 4 & -1 & -1 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline 1 & 3 & 2 & 1 \\ \hline 1 & 0 & 0 & -2 \\ \hline 4 & -2 & 3 & 0 \\ \hline 0 & 1 & 2 & -5 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 19 & -3 \\ \hline 2 & 3 \\ \hline \end{array}$$

Convolutional neural networks (CNNs)

- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:**

$$\begin{array}{|c|c|c|} \hline 1 & 3 & -1 \\ \hline 0 & -2 & 0 \\ \hline 4 & -1 & -1 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline 1 & 3 & 2 & 1 \\ \hline 1 & 0 & 0 & -2 \\ \hline 4 & -2 & 3 & 0 \\ \hline 0 & 1 & 2 & -5 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 19 & -3 \\ \hline 2 & 3 \\ \hline \end{array}$$

Convolutional neural networks (CNNs)

- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:**

$$\begin{array}{|c|c|c|} \hline 1 & 3 & -1 \\ \hline 0 & -2 & 0 \\ \hline 4 & -1 & -1 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline 1 & 3 & 2 & 1 \\ \hline 1 & 0 & 0 & -2 \\ \hline 4 & -2 & 3 & 0 \\ \hline 0 & 1 & 2 & -5 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 19 & -3 \\ \hline 2 & 3 \\ \hline \end{array}$$

Convolutional neural networks (CNNs)

- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:** keep dimension by **padding** with zeroes

$$\begin{array}{|c|c|c|} \hline 1 & 3 & -1 \\ \hline 0 & -2 & 0 \\ \hline 4 & -1 & -1 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline 1 & 3 & 2 & 1 \\ \hline 1 & 0 & 0 & -2 \\ \hline 4 & -2 & 3 & 0 \\ \hline 0 & 1 & 2 & -5 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 19 & -3 \\ \hline 2 & 3 \\ \hline \end{array}$$

Convolutional neural networks (CNNs)

- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:** keep dimension by **padding** with zeroes

$$\begin{array}{|c|c|c|} \hline 1 & 3 & -1 \\ \hline 0 & -2 & 0 \\ \hline 4 & -1 & -1 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline 1 & 3 & 2 & 1 \\ \hline 1 & 0 & 0 & -2 \\ \hline 4 & -2 & 3 & 0 \\ \hline 0 & 1 & 2 & -5 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 19 & -3 \\ \hline 2 & 3 \\ \hline \end{array}$$

- **Generalization of matrix multiplication:** tensors can have higher dimensions

Convolutional neural networks (CNNs)

- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:** keep dimension by **padding** with zeroes

1	3	-1
0	-2	0
4	-1	-1

 *

1	3	2	1
1	0	0	-2
4	-2	3	0
0	1	2	-5

 =

19	-3
2	3

- **Generalization of matrix multiplication:** tensors can have higher dimensions

The diagram shows a 3D tensor of size $D \times D \times c$ (represented as a rectangular prism with a dashed square on the top face) multiplied by a 3D tensor of size $d \times d \times c$ (represented as a smaller rectangular prism with a dashed square on the top face). The result is a 2D tensor of size $D \times D$ (represented as a square).

Convolutional neural networks (CNNs)


- **Generality:** we actually prove the result in CNNs
- **Discrete convolution:** keep dimension by **padding** with zeroes

$$\begin{bmatrix} 1 & 3 & -1 \\ 0 & -2 & 0 \\ 4 & -1 & -1 \end{bmatrix} * \begin{bmatrix} 1 & 3 & 2 & 1 \\ 1 & 0 & 0 & -2 \\ 4 & -2 & 3 & 0 \\ 0 & 1 & 2 & -5 \end{bmatrix} = \begin{bmatrix} 19 & -3 \\ 2 & 3 \end{bmatrix}$$

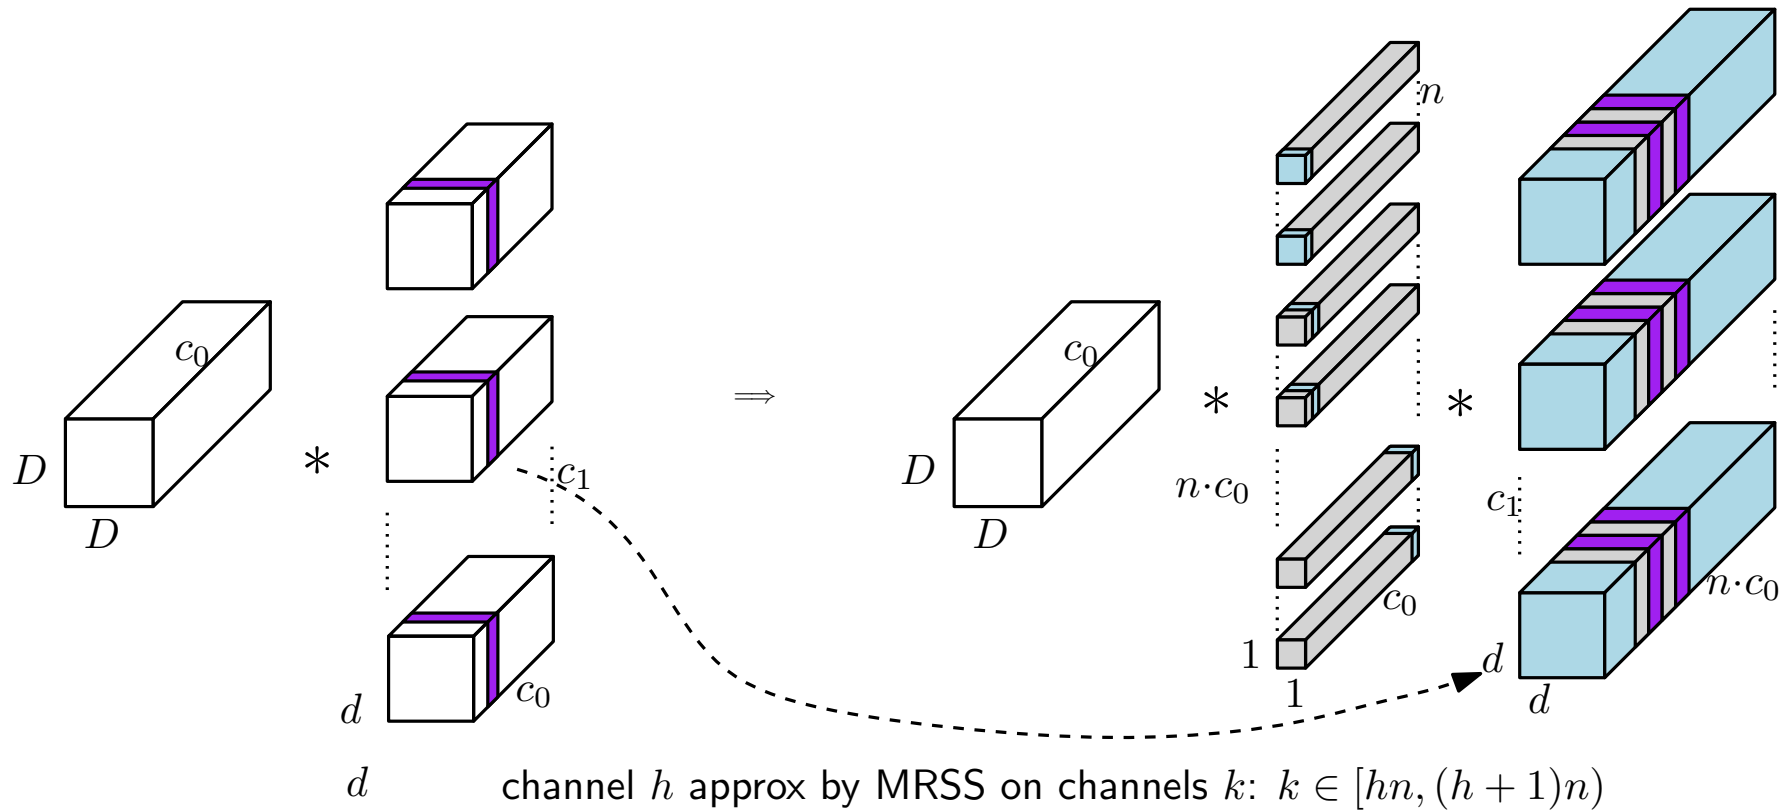
- **Generalization of matrix multiplication:** tensors can have higher dimensions

The diagram shows a 3D tensor of size $D \times D \times c$ multiplied by a stack of n 3D tensors, each of size $d \times d \times c$. The result is a 3D tensor of size $D \times D \times n$.

SLTH construction in CNNs




SLTH construction in CNNs




- **Restrictions** on the structure of the CNN

SLTH construction in CNNs



- **Restrictions** on the structure of the CNN
- Only **ReLU** activation function

SLTH construction in CNNs



- **Restrictions** on the structure of the CNN
- Only **ReLU** activation function
- $n \geq \text{poly}(d) \cdot \text{polylog}(d\ell/\varepsilon)$ is sufficient

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**
- **Solution:** multidimensional RSS

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**
- **Solution:** multidimensional RSS
- **Modifications:** adaptation of MRSS to random vectors with **dependent entries**

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**
- **Solution:** multidimensional RSS
- **Modifications:** adaptation of MRSS to random vectors with **dependent entries**
- **Our result:** SLTH holds in CNNs via **structured pruning** with **polynomial overparameterization**

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**
- **Solution:** multidimensional RSS
- **Modifications:** adaptation of MRSS to random vectors with **dependent entries**
- **Our result:** SLTH holds in CNNs via **structured pruning** with **polynomial overparameterization**
- **Open (1):** tightness of MRSS ($n \geq d \log 1/\varepsilon$?)

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**
- **Solution:** multidimensional RSS
- **Modifications:** adaptation of MRSS to random vectors with **dependent entries**
- **Our result:** SLTH holds in CNNs via **structured pruning** with **polynomial overparameterization**
- **Open (1):** tightness of MRSS ($n \geq d \log 1/\varepsilon$?)
- **Open (2):** how to replace the union bound?

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**
- **Solution:** multidimensional RSS
- **Modifications:** adaptation of MRSS to random vectors with **dependent entries**
- **Our result:** SLTH holds in CNNs via **structured pruning** with **polynomial overparameterization**
- **Open (1):** tightness of MRSS ($n \geq d \log 1/\varepsilon$?)
- **Open (2):** how to replace the union bound?
- **Open (3):** generalization of CNN structure, activation function, etc.

Conclusions

- **Previously:** SLTH holds via **unstructured pruning** in dense networks, CNNs, etc., with logarithmic overhead
- **Tool:** the one-dimensional RSS problem is heavily exploited
- **Issue:** it leads to **exponential bounds** when trying to achieve **structured pruning**
- **Solution:** multidimensional RSS
- **Modifications:** adaptation of MRSS to random vectors with **dependent entries**
- **Our result:** SLTH holds in CNNs via **structured pruning** with **polynomial overparameterization**
- **Open (1):** tightness of MRSS ($n \geq d \log 1/\varepsilon$?)
- **Open (2):** how to replace the union bound?
- **Open (3):** generalization of CNN structure, activation function, etc.

Thank you!

RSS proof overview


- [Lueker 1998; da Cunha et al. 2023]

RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$




RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$




RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$




RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$




RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$




RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$




RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$
- Approximate the **whole interval**




RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$
- Approximate the **whole interval**




Consider
$$f_t(x) = \begin{cases} 1 & \text{if } x \in [-1, 1] \text{ and } \exists S \subseteq [t] : |x - \sum_{i \in S} X_i| < 2\varepsilon \\ 0 & \text{otherwise} \end{cases}$$

RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$
- Approximate the **whole interval**



Consider $f_t(x) = \begin{cases} 1 & \text{if } x \in [-1, 1] \text{ and } \exists S \subseteq [t] : |x - \sum_{i \in S} X_i| < 2\varepsilon \\ 0 & \text{otherwise} \end{cases}$


$v_t = \frac{1}{2} \int_{-1}^1 f_t(x) dx$ keeps track of the **approximated volume**

RSS proof overview

- [Lueker 1998; da Cunha et al. 2023]

Specific instance of RSSP

- X_1, \dots, X_n **uniform** random variables over $[-1, 1]$
- **Error** parameter $\varepsilon > 0$
- Approximate the **whole interval**



Consider $f_t(x) = \begin{cases} 1 & \text{if } x \in [-1, 1] \text{ and } \exists S \subseteq [t] : |x - \sum_{i \in S} X_i| < 2\varepsilon \\ 0 & \text{otherwise} \end{cases}$

$v_t = \frac{1}{2} \int_{-1}^1 f_t(x) dx$ keeps track of the **approximated volume**

By restricting f_t , v_t becomes a sub-martingale